

Security Monitoring in the Cloud
Session #3

Application & Data Security Monitoring

Jan 2026

Key Outcomes

- Implement observability for cloud applications using logs and traces.
- Detect unauthorized access to cloud storage and data services.
- Create actionable alerts for anomalous data activity.
- Integrate monitoring outputs with SIEM workflows.

Course Material: https://github.com/0x74696D/security_monitoring_tp

Agenda

01 **Observability for Cloud based Applications**
Logging patterns and integrations

02 **Monitoring Data Services**
Storage, DBs, Queues, etc.

03 **Alerting**
Best practices and threat scenarios

Observability for Cloud Applications

Observability for Cloud Applications

1.0.0

Cloud Logging for Applications

Security logs:



- Auth events (login success/fail)



- Permission changes, role changes



- Sensitive operations (exports, deletes, config changes)



- Client IP, user ID, request IDs

GCP Cloud Logging or AWS CloudWatch Logs for:

- Serverless workloads (Cloud Run Functions/Jobs, Lambda, Batch, Step Functions, etc.)
 - Containerized workloads (Cloud Run Services, ECS)
 - Kubernetes workloads (GKE, EKS)
 - Load balancers, API Gateway
-
- + CDN logs (Cloud HTTP(S) Load Balancer logs for GCP or Access Logs for Cloudfront)
 - + Firewall logs (Google Cloud Armor, WAF)



Observability for Cloud Applications

1.1.0

Logging Patterns



Who did what, from where, to which resource, and was it allowed.

Question	Example
Who	user_id, service_account
What	action (login, export, delete, grant_role)
Where	source_ip, region, user_agent
When	timestamp (always!)
On what	resource_id (file, record, dataset)
Result	success / failure / denied
Context	request_id / trace_id

Observability for Cloud Applications

1.2.0

Best Practices



Structured Logging

- machine-readable, not free-text blobs
- standardized fields across systems
- security-relevant events, not raw noise: high-value events



```
{  
  "timestamp": "2026-01-24T10:31:42Z",  
  "event_type": "auth_failure",  
  "user_id": "12345",  
  "source_ip": "203.0.113.42",  
  "service": "api-gateway",  
  "destination": "auth_service_prod",  
  "reason": "invalid_password",  
  "environment": "prod",  
  "request_id": "019bf076-916f-7d10-88c4-f776c9bd0587"
```

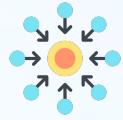
User failed to login from 203.0.113.42 at
10:31:42



Observability for Cloud Applications

1.2.1

Best Practices



Centralized

- not live only on the host
- RBAC, encryption , monitoring
- retention policies: hot/cold/archive
- **Correlation**



Alert on behavior, not single events

- 20 failed logins in 5 minutes
- Login from new country + privilege escalation
- API key used from two regions simultaneously



Correlate logs with metrics & traces

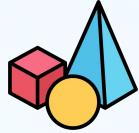
- Auth failures + traffic spikes
- WAF blocks + application errors
- IAM changes + unusual API usage

Monitoring Data Services

Monitoring Data Services

2.0.0

Monitoring blob access



GCS access logs or AWS S3 access logs & audit logs or CloudTrail:

`storage.objects.get, storage.objects.list, storage.objects.delete`

What to monitor:



- Bulk downloads
- Access from unusual identities / service accounts
- Access from unexpected regions or IPs

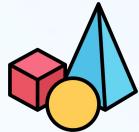
Example scenarios:



- Exfiltration from a “backups” bucket
- Public bucket misconfig → mass access

Monitoring Data Services

Monitoring Databases & Data Services



Common data backends:

- SQL databases (PaaS) Cloud SQL / RDS, Aurora
- Data Warehouse BigQuery / Redshift
- No SQL DB: Firestore, DynamoDB



What's important to log:

- Login attempts (DB auth failures)
- Privilege changes (GRANT/REVOKE)
- Mass SELECTs, EXPORTs, COPY TO statements

Monitoring Data Services

Monitoring Databases & Data Services

2.2.0

Data Service	Cloud Provider	Logging Service	What Gets Logged (Security-Relevant)
Cloud SQL	GCP	Cloud Audit Logs + DB engine logs	Logins, failed auth, schema changes, connections
BigQuery	GCP	Cloud Audit Logs (Data Access)	Queries, extracts, exports, dataset access
Cloud Storage (GCS)	GCP	Cloud Audit Logs (Data Access)	Object reads, writes, deletes, bulk downloads
Firestore / Datastore	GCP	Cloud Audit Logs	Document reads/writes, admin operations
Aurora (MySQL/Postgres)	AWS	CloudWatch Logs	SQL queries, connections, auth failures
RDS (MySQL/Postgres)	AWS	CloudWatch Logs	Login attempts, slow queries, errors
S3	AWS	CloudTrail + S3 Access Logs	Object access, list/get/delete operations
DynamoDB	AWS	CloudTrail + CloudWatch Metrics	Table access, scan/query operations

Alerting

Alerting

3.0.0

From Observability to Detection (What Becomes an Alert?)

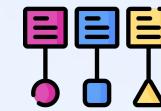


Not every log is an alert → building *signal* from noise



- **What events should be alerting:**

- Login failures threshold (brute force)
- Privilege escalation actions
- Bulk exports
- Anomalous regions or IPs



- **Alert types:**

- Threshold-based (N events in T minutes)
- Pattern-based (specific queries or endpoints)
- Behavior-based (deviation from baseline)

Alerting

Best Practices for Alert Design

3.1.0



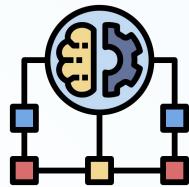
Avoiding alert fatigue

- Prioritize “actionable” over “perfect coverage”
- Use severity levels (info / low / med / high / critical)



Designing for investigation

- Include context (user, IP, trace ID, resource)
- Include runbook links / recommended actions



Map to Frameworks

- MITRE stages (recon, access, exfil)
- Incident response workflows

Alerting

Alert Design Steps

Detect malicious behavior, not isolated events.

Attacker Behavior

Required Telemetry

Atomic Detections

Correlation in Time & Context

Actionability



Step 1 - Start from Attacker Behavior

What is the attacker trying to do?

- Access credentials?
- Escalate privileges?
- Move laterally?
- Exfiltrate data?

→ Translate behavior into **observable signals**

Alerting

Alert Design Steps

Detect malicious behavior, not isolated events.

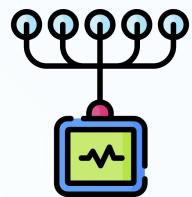
Attacker Behavior

Required Telemetry

Atomic Detections

Correlation in Time & Context

Actionability



Step 2 - Identify Required Telemetry

Map each behavior to logs

Behavior	Telemetry
Identity abuse	IAM / API audit logs
Cluster abuse	Kubernetes audit logs
Runtime execution	Container / node logs
Lateral movement	Network flow logs
Data access	Storage / DB audit logs

Alerting

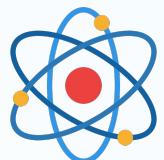
Alert Design Steps

Detect malicious behavior, not isolated events.



Step 3 - Build Atomic Detections

Create small, reliable signals:



- “Service account used from new IP”
- “Pod exec in prod namespace”
- “Privileged pod created”
- “Bulk data download”

These should be **low-noise, high-confidence**.

Alerting

Alert Design Steps

Detect malicious behavior, not isolated events.



Step 4 - Correlate in Time & Context



- *Link signals across layers*
 - *Same identity*
 - *Same workload*
 - *Same time window*
 - *Same resource*

Alerting

Alert Design Steps

Detect malicious behavior, not isolated events.



Step 5 - Tune for Actionability



- *Suppress known automation*
- *Add thresholds & time windows*
- *Attach investigation context*
- *Map to MITRE & response playbooks*

Question ?

Additional Resources

1. [OWASP Logging Best Practices](#)
2. [Guide to Computer Security Log Management](#)

Hands-on time after a small break !



Detect reconnaissance activity within GKE

https://github.com/0x74696D/security_monitoring_tp